

Warning: Do not run notebook, certain cells take upwards of 30 minutes to an hour to run!!!

Petfinder Pawpularity Prediction Using Neural Networks

Authors: Aisha Baitemirova-Othman, Angela Kim, Steven Addison

Contents

- [Overview](#)
 - [Business Understanding](#)
 - [Data Understanding](#)
 - [Imports](#)
 - [Data Preparation & Analysis](#)
 - [Modeling](#)
 - [Baseline Dummy Regressor](#)
 - [First Basic Model and Image Augmentation](#)
 - [Iterative Artificial Neural Networks](#)
 - [Iterative Convolutional Neural Networks](#)
 - [Conclusion](#)
 - [Next Steps](#)
 - [Sources](#)
-

Overview

[Petfinder.my \(https://www.petfinder.my/\)](https://www.petfinder.my/) is Malaysia's leading animal welfare platform, featuring over 200,000 animals with more than 56,000 happily adopted. This project analyzes photos of adoptable pets from Malaysian animal shelters found on Petfinder and designs a Deep Learning model to predict the "Pawpularity" of pet photos.

Business Understanding

They say a picture is worth a thousand words. A picture can also save a life. Hundreds of millions of stray cats and dogs suffer on the streets, live miserably in crowded shelters, or are euthanized around the world. Companion animals with attractive and high quality photos are more likely to be adopted and more likely to be adopted faster [source](#) (<https://www.tandfonline.com/doi/full/10.1080/10888705.2014.982796>). We want to answer the question: what makes a good picture? After analyzing raw images and metadata to predict the “Pawpularity” of pet photos, we train and test our model on PetFinder.my's thousands of pet profiles to come up with the best recommendations on photo composition. We hope our model will help stray cats and dogs find their "forever" home faster.

Data Understanding

The data comes from thousands of pet profiles on [Petfinder.my](#) (<https://www.petfinder.my/>). The Pawpularity score is derived from each pet profile's page view statistics at the listing pages, using an algorithm that normalizes the traffic data across different pages, platforms (web & mobile) and various metrics. Duplicate clicks, crawler bot accesses and sponsored profiles are excluded from the analysis. All the feature metadata is explained [here](#) (<https://github.com/stevenaddison/Project-4/blob/main/data/metadata.md>).

Imports

First, we start with importing the relevant libraries to load and clean our dataset.

```
In [1]: #pip install opencv-python  
#pip install scikit-image
```

```
In [2]: import numpy as np
import pandas as pd
import cv2
import tensorflow as tf
import keras

from skimage.transform import resize
from skimage import color, io

from scipy import ndimage

from sklearn.model_selection import train_test_split
from sklearn.dummy import DummyRegressor
from sklearn.metrics import mean_squared_error as mse

from keras.models import Sequential
from keras.layers import Dense, Flatten, Conv2D, MaxPooling2D, Dropout
from keras.losses import MeanSquaredError
from keras.metrics import RootMeanSquaredError

import matplotlib.pyplot as plt
import seaborn as sns

import warnings
warnings.filterwarnings('ignore')
```

/opt/anaconda3/envs/learn-env/lib/python3.8/site-packages/skimage/io/manage_plugins.py:23: UserWarning: Your installed pillow version is < 8.1.2. Several security issues (CVE-2021-27921, CVE-2021-25290, CVE-2021-25291, CVE-2021-25293, and more) have been fixed in pillow 8.1.2 or higher. We recommend to upgrade this library.

```
from .collection import imread_collection_wrapper
```

Data Preparation & Analysis

```
In [3]: # Open .csv file as pandas dataframe
df = pd.read_csv("data/train.csv")
```

```
In [4]: # Initial Look at dataframe
df.head()
```

Out[4]:

		Id	Subject Focus	Eyes	Face	Near	Action	Accessory	Group	Coll
0	0007de18844b0dbbb5e1f607da0606e0		0	1	1	1	0	0	1	
1	0009c66b9439883ba2750fb825e1d7db		0	1	1	0	0	0	0	
2	0013fd999caf9a3efe1352ca1b0d937e		0	1	1	1	0	0	0	
3	0018df346ac9c1d8413cfcc888ca8246		0	1	1	1	0	0	0	
4	001dc955e10590d3ca4673f034feef2		0	0	0	1	0	0	1	

In [5]: *# Count of non-null values, datatypes, and total entries*
df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 9912 entries, 0 to 9911
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Id                    9912 non-null   object
1   Subject Focus         9912 non-null   int64
2   Eyes                  9912 non-null   int64
3   Face                  9912 non-null   int64
4   Near                  9912 non-null   int64
5   Action                9912 non-null   int64
6   Accessory             9912 non-null   int64
7   Group                 9912 non-null   int64
8   Collage               9912 non-null   int64
9   Human                 9912 non-null   int64
10  Occlusion              9912 non-null   int64
11  Info                  9912 non-null   int64
12  Blur                  9912 non-null   int64
13  Pawpularity           9912 non-null   int64
dtypes: int64(13), object(1)
memory usage: 1.1+ MB
```

In [6]: *# Check descriptive statistics*
df.describe()

Out[6]:

	Subject Focus	Eyes	Face	Near	Action	Accessory	Group
count	9912.000000	9912.000000	9912.000000	9912.000000	9912.000000	9912.000000	9912.000000
mean	0.027643	0.772599	0.903955	0.861582	0.009988	0.067797	0.129338
std	0.163957	0.419175	0.294668	0.345356	0.099444	0.251409	0.335591
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
25%	0.000000	1.000000	1.000000	1.000000	0.000000	0.000000	0.000000
50%	0.000000	1.000000	1.000000	1.000000	0.000000	0.000000	0.000000
75%	0.000000	1.000000	1.000000	1.000000	0.000000	0.000000	0.000000
max	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000

```
In [7]: # Check value counts
for c in df.columns:
    print ("---- %s ----" % c)
    print (df[c].value_counts())
    print("\n")
```

```
---- Id ----
a9894bf4223185f48188a545044aeefa    1
c97498a8f67676fed254ae35e735a312    1
d20cf059a9e892dcb8bb991bf5acf010    1
b6004a1d03498ef6ef362685e38535e2    1
2aee7c914c13f432e20c842ce71dbfd9    1
..
a450cd16e284e1fc9affba674a63c154    1
4a4b92b7387e75223b2fd1e502e71ce3    1
5dadbfa60882b172cc2f0710aaefb9f1    1
87abac899ef38e21df16edab73e05307    1
efc51b5700c80e739d81f83bd21ef735    1
Name: Id, Length: 9912, dtype: int64
```

```
---- Subject Focus ----
0    9638
1     274
Name: Subject Focus, dtype: int64
```

```
---- Eyes ----
1    7658
0    2254
Name: Eyes, dtype: int64
```

```
---- Face ----
1    8960
0     952
Name: Face, dtype: int64
```

```
---- Near ----
1    8540
0    1372
Name: Near, dtype: int64
```

```
---- Action ----
0    9813
1      99
Name: Action, dtype: int64
```

```
---- Accessory ----
0    9240
1     672
Name: Accessory, dtype: int64
```

```

---- Group ----
0      8630
1      1282
Name: Group, dtype: int64

```

```

---- Collage ----
0      9420
1       492
Name: Collage, dtype: int64

```

```

---- Human ----
0      8264
1      1648
Name: Human, dtype: int64

```

```

---- Occlusion ----
0      8207
1      1705
Name: Occlusion, dtype: int64

```

```

---- Info ----
0      9305
1       607
Name: Info, dtype: int64

```

```

---- Blur ----
0      9214
1       698
Name: Blur, dtype: int64

```

```

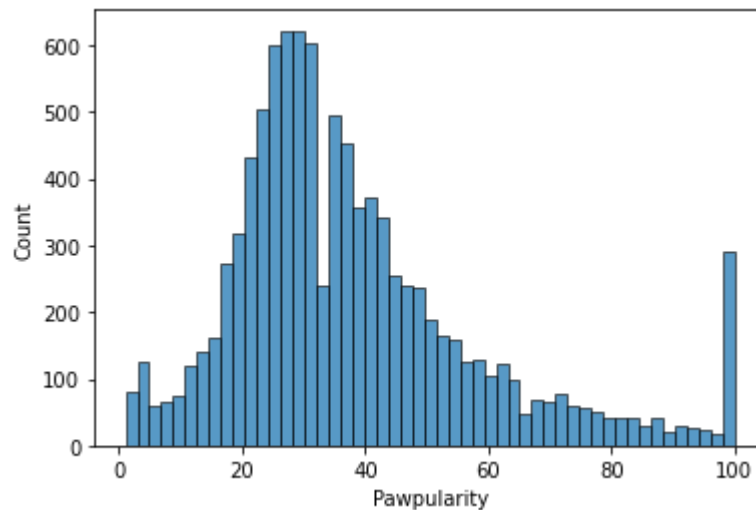
---- Pawpularity ----
28      318
30      318
26      316
31      312
29      304
...
98       10
97        8
90        7
1         4
99        4
Name: Pawpularity, Length: 100, dtype: int64

```

There are 9912 images and 12 features: Subject Focus , Eyes , Face , Near , Action , Accessory , Group , Collage , Human , Occlusion , Info , and

Blur . The target variable is `Pawpularity` and ranges from 1-100. There are no null values and all features have a value of 0 (no) or 1 (yes).

In [8]: `sns.histplot(df['Pawpularity']);`



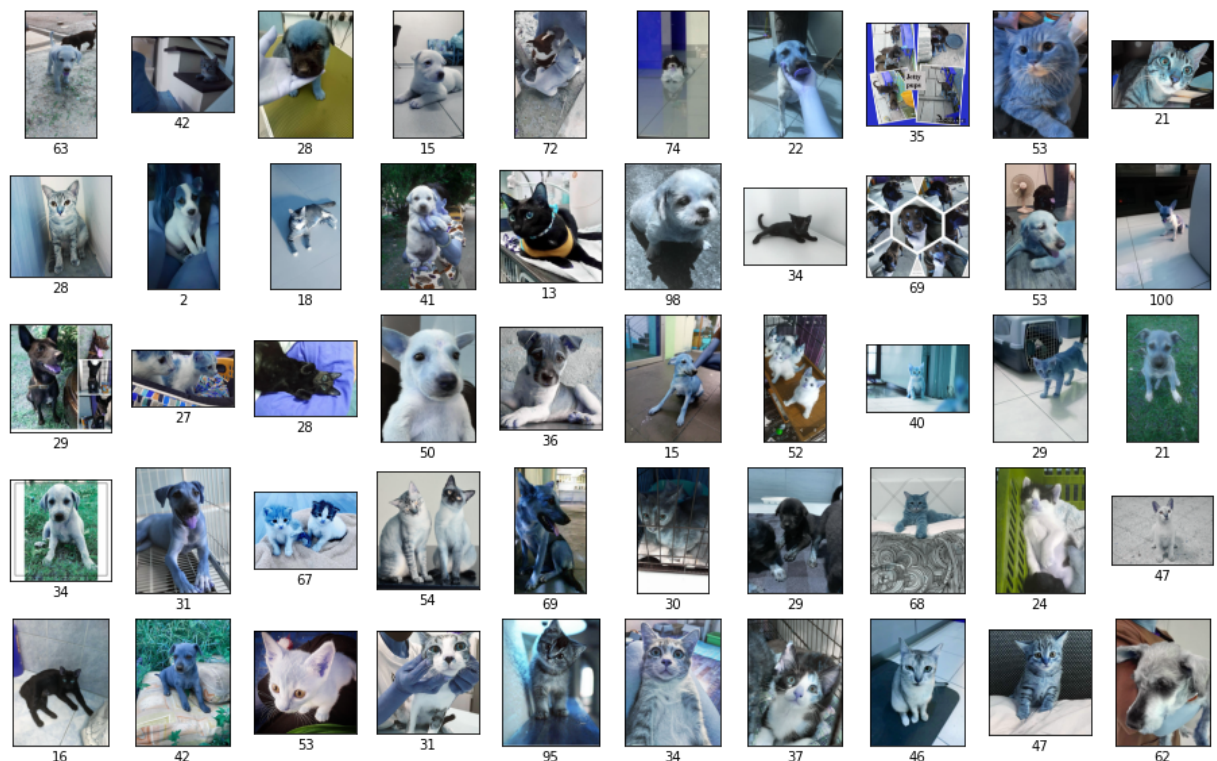
Here we take a look at the distribution of the target `Pawpularity` and see a spike at 100.

```
In [9]: # Download images to notebook
images = []
count = 0
for index in range(len(df)):
    Id = df.loc[index, 'Id']
    path = 'data/train/'+str(Id)+'.jpg'
    img_array = cv2.imread(path)
    img_array = resize(img_array, (128, 128), anti_aliasing=True)
    images.append(img_array)
    count += 1
    progress = (count/len(df))*100
    print('progress =', round(progress,2), '%', end='\r')
```

progress = 100.0 %12.08 %

```
In [10]: # Plot first 50 images
def image_read(path):
    return cv2.imread(path)

df['path'] = df.apply(lambda x : 'data/train/' + x['Id'] + ".jpg", axis=1)
fig, ax = plt.subplots(5,10,figsize=(16, 10))
for i, (path, score) in enumerate(df[['path', 'Pawpularity']][:50].values.tolist()):
    row, col = i // 10, i % 10
    axis = ax[row][col]
    axis.imshow(image_read(path))
    axis.set_xticks([])
    axis.set_yticks([])
    axis.set_xlabel(score)
plt.show()
```



We see a variety of cat and dog photos and their Pawpularity score. Something we notice is that the photos lean blue but do not know why and cannot find more information on this.


```
In [11]: # Transforming images list to numpy array for modeling and assigning it to X and
X = np.array(images)
y = df['Pawpularity']

# Checking shape of X
X.shape
```

```
Out[11]: (9912, 128, 128, 3)
```

```
In [12]: # Do a train-test-split using a random state of 42 for reproducibility
X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=42)
```

Modeling

We start by creating functions for model building and evaluation to streamline our process.

```
In [13]: # Function for model building
def model_compfitt(model, X_train, y_train, n_epochs):
    """Function to compile our models using our chosen optimizer, loss,
    and evaluation metrics. The function then fits the model on our train
    data and n number epochs, finally printing out a model summary."""

    # compile model with adam optimizer and mean squared error metric
    model.compile(optimizer='adam',
                  loss=MeanSquaredError(),
                  metrics=[RootMeanSquaredError()])
    # fit model on train sets with n number epochs
    model.fit(X_train, y_train, epochs=n_epochs)

    #print model summary
    model.summary()
```

```

In [14]: # Function for model evaluation
def model_eval(model, X_train, y_train, X_test, y_test, name):
    """Function to evaluate our model on the training and test
    data, then printing out those metrics, and visualizing the
    predictions on the test data."""

    # calculate training data metrics
    train_eval = model.evaluate(X_train, y_train)

    # calculate test data metrics
    test_eval = model.evaluate(X_test, y_test)

    # print metrics
    print(f"""
    {name} Training Metrics:
    Loss: {round(train_eval[0] , 3)}
    RMSE: {round(train_eval[1] , 3)}
    -----
    {name} Test Metrics:
    Loss: {round(test_eval[0] , 3)}
    RMSE: {round(test_eval[1] , 3)}
    """)

    # visualize predictions
    ypred = model.predict(X_test)
    x_ax = range(len(ypred))
    plt.figure(figsize=(20,10))
    plt.scatter(x_ax, y_test, s=15, color="blue", label="original")
    plt.plot(x_ax, ypred, lw=1, color="red", label="predicted")
    plt.legend()
    plt.show()

```

Baseline Dummy Regressor

We start off with a Baseline Dummy Regressor to compare our models to and try to improve upon.

```

In [15]: baseline = DummyRegressor(strategy="mean")
baseline.fit (X_train,y_train)
y_hat_test = baseline.predict(X_test)
baseline_rmse = mse(y_test, y_hat_test, squared=False)

```

```

In [16]: baseline_rmse

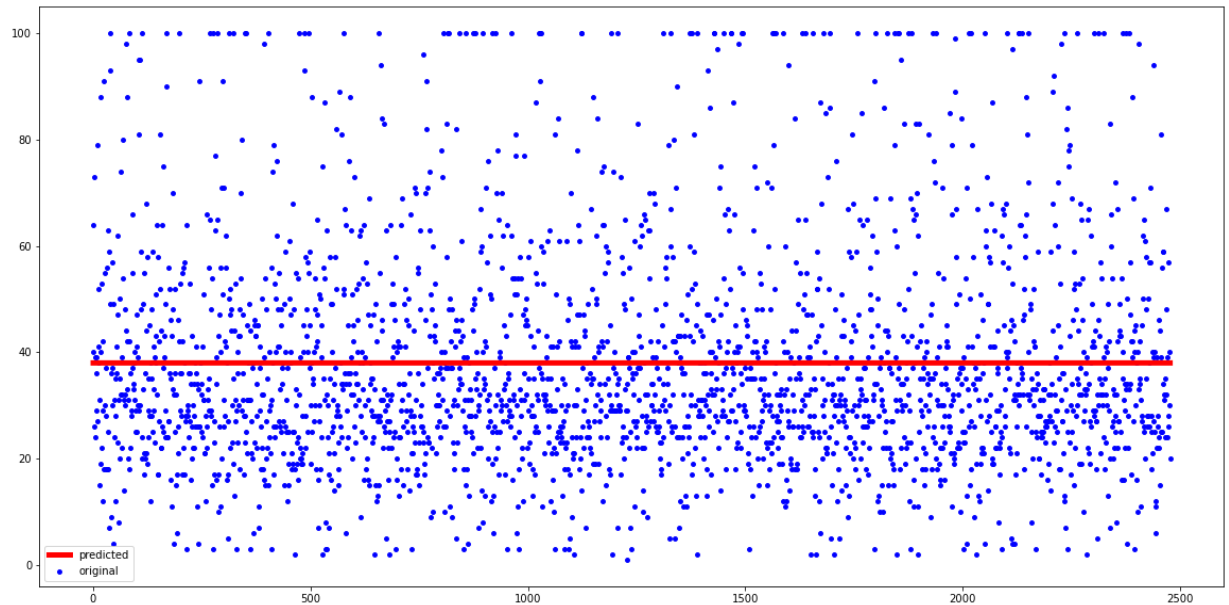
```

```

Out[16]: 21.074920522735773

```

```
In [17]: ypred = baseline.predict(X_test)
x_ax = range(len(ypred))
plt.figure(figsize=(20,10))
plt.scatter(x_ax, y_test, s=15, color="blue", label="original")
plt.plot(x_ax, ypred, lw=5, color="red", label="predicted")
plt.legend()
plt.show()
```



First Basic Model and Image Augmentation

Our first basic model is an Artificial Neural Network with one layer.

```
In [18]: basicann = Sequential([
          Flatten(),
          Dense(1, activation='relu')
        ])

model_compfit(basicann, X_train, y_train, 10)
```

```
Epoch 1/10
233/233 [=====] - 2s 7ms/step - loss: 498.5932 - root_
mean_squared_error: 22.3292
Epoch 2/10
233/233 [=====] - 1s 5ms/step - loss: 499.8664 - root_
mean_squared_error: 22.3577
Epoch 3/10
233/233 [=====] - 1s 5ms/step - loss: 479.0622 - root_
mean_squared_error: 21.8875
Epoch 4/10
233/233 [=====] - 1s 5ms/step - loss: 488.3990 - root_
mean_squared_error: 22.0998
Epoch 5/10
233/233 [=====] - 1s 4ms/step - loss: 474.0719 - root_
mean_squared_error: 21.7732
Epoch 6/10
233/233 [=====] - 1s 5ms/step - loss: 471.8141 - root_
mean_squared_error: 21.7213
Epoch 7/10
233/233 [=====] - 1s 5ms/step - loss: 471.9232 - root_
mean_squared_error: 21.7238
Epoch 8/10
233/233 [=====] - 2s 10ms/step - loss: 465.8413 - root_
_mean_squared_error: 21.5834
Epoch 9/10
233/233 [=====] - 2s 8ms/step - loss: 458.4019 - root_
mean_squared_error: 21.4103
Epoch 10/10
233/233 [=====] - 1s 5ms/step - loss: 470.1948 - root_
mean_squared_error: 21.6840
Model: "sequential"
```

Layer (type)	Output Shape	Param #
flatten (Flatten)	(None, 49152)	0
dense (Dense)	(None, 1)	49153

Total params: 49,153
 Trainable params: 49,153
 Non-trainable params: 0

```
In [19]: model_eval(basicann, X_train, y_train, X_test, y_test, 'Basic ANN')
```

```
233/233 [=====] - 2s 10ms/step - loss: 453.5609 - root_mean_squared_error: 21.2970
78/78 [=====] - 0s 5ms/step - loss: 528.6611 - root_mean_squared_error: 22.9926
```

Basic ANN Training Metrics:

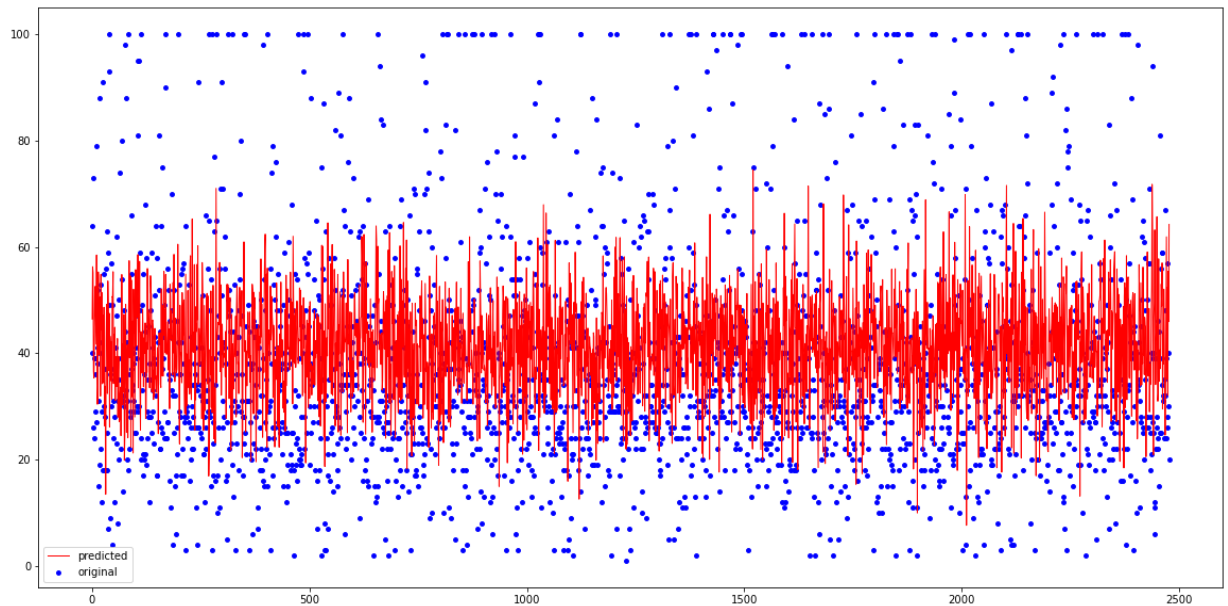
Loss: 453.561

RMSE: 21.297

Basic ANN Test Metrics:

Loss: 528.661

RMSE: 22.993



We do some image augmentation by blurring the photos to see how that affects our RMSE.

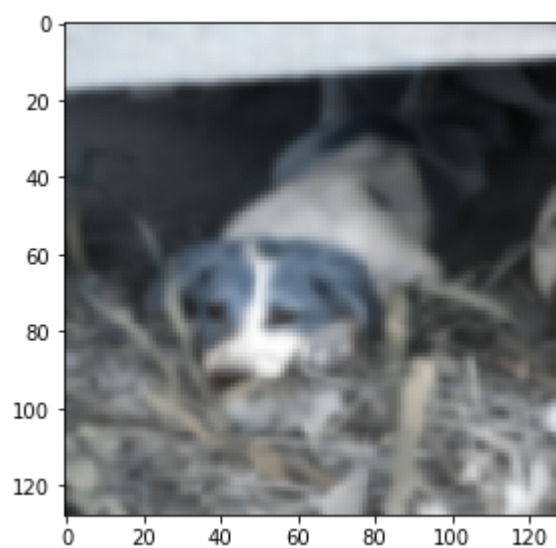
```
In [20]: X_train_blur = []
```

```
for image in X_train:
    image = ndimage.gaussian_filter(image, sigma= 1)
    X_train_blur.append(image)

X_train_blur = np.array(X_train_blur)
```

```
In [21]: io.imshow(X_train_blur[0])
```

```
Out[21]: <matplotlib.image.AxesImage at 0x7fabae4590d0>
```



```
In [22]: blurann = Sequential([
            Flatten(),
            Dense(1, activation='relu')
        ])

model_compfit(blurann, X_train_blur, y_train, 10)
```

```
Epoch 1/10
233/233 [=====] - 1s 4ms/step - loss: 508.3556 - root_
mean_squared_error: 22.5467
Epoch 2/10
233/233 [=====] - 1s 4ms/step - loss: 488.4328 - root_
mean_squared_error: 22.1005
Epoch 3/10
233/233 [=====] - 1s 4ms/step - loss: 498.3028 - root_
mean_squared_error: 22.3227
Epoch 4/10
233/233 [=====] - 1s 4ms/step - loss: 491.4557 - root_
mean_squared_error: 22.1688
Epoch 5/10
233/233 [=====] - 1s 4ms/step - loss: 489.2918 - root_
mean_squared_error: 22.1199
Epoch 6/10
233/233 [=====] - 1s 4ms/step - loss: 478.0587 - root_
mean_squared_error: 21.8646
Epoch 7/10
233/233 [=====] - 1s 4ms/step - loss: 487.8592 - root_
mean_squared_error: 22.0875
Epoch 8/10
233/233 [=====] - 1s 4ms/step - loss: 482.4075 - root_
mean_squared_error: 21.9638
Epoch 9/10
233/233 [=====] - 1s 5ms/step - loss: 483.1240 - root_
mean_squared_error: 21.9801
Epoch 10/10
233/233 [=====] - 2s 9ms/step - loss: 474.9717 - root_
mean_squared_error: 21.7938
Model: "sequential_1"
```

Layer (type)	Output Shape	Param #
flatten_1 (Flatten)	(None, 49152)	0
dense_1 (Dense)	(None, 1)	49153

Total params: 49,153
 Trainable params: 49,153
 Non-trainable params: 0

```
In [23]: model_eval(blurann, X_train_blur, y_train, X_test, y_test, 'Blurred ANN')
```

```
233/233 [=====] - 4s 19ms/step - loss: 470.5186 - root_mean_squared_error: 21.6914
78/78 [=====] - 1s 8ms/step - loss: 541.4640 - root_mean_squared_error: 23.2694
```

Blurred ANN Training Metrics:

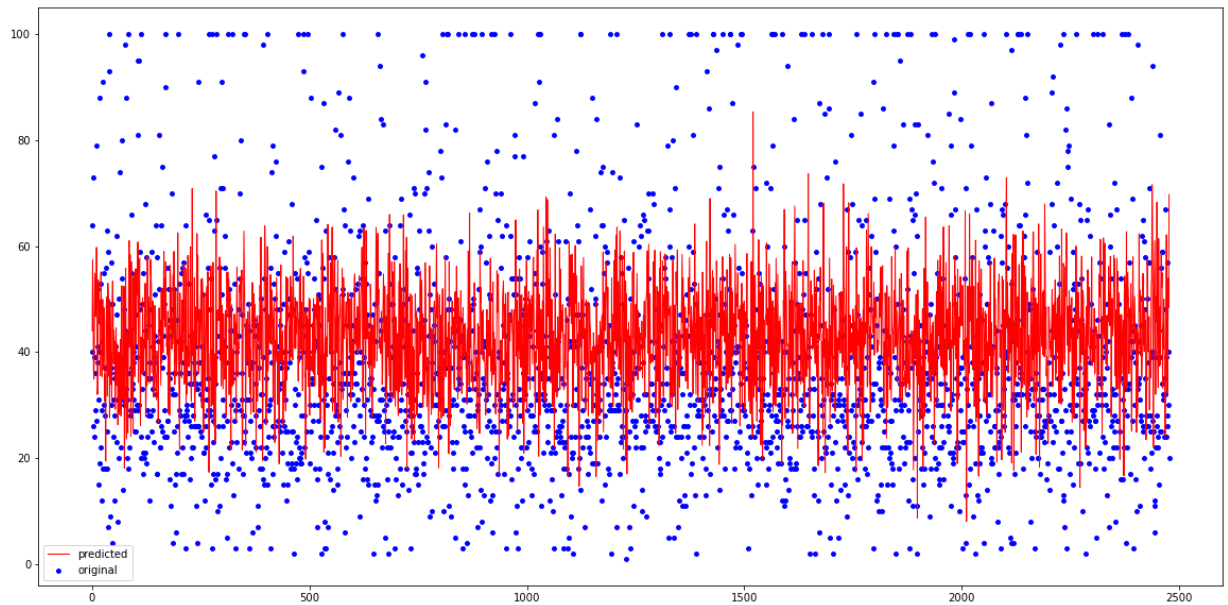
Loss: 470.519

RMSE: 21.691

Blurred ANN Test Metrics:

Loss: 541.464

RMSE: 23.269



We then flip the images horizontally.

```
In [24]: X_train_flip = []

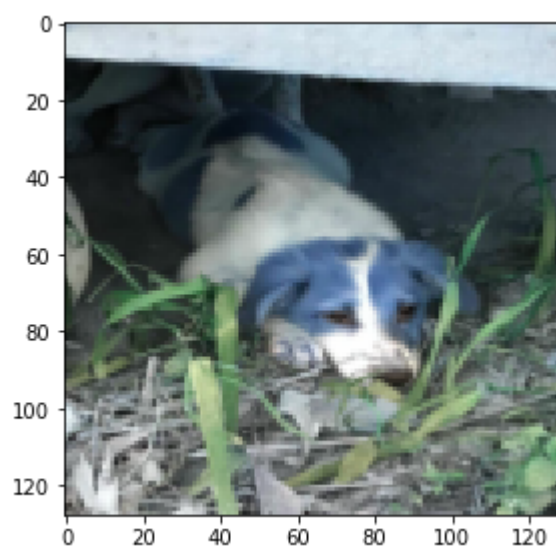
for image in X_train:
    image = cv2.flip(image,1)
    X_train_flip.append(image)

X_train_flip = np.array(X_train_flip)
```



```
In [25]: io.imshow(X_train_flip[0])
```

```
Out[25]: <matplotlib.image.AxesImage at 0x7faa9d8b3a90>
```



```
In [26]: flipann = Sequential([
            Flatten(),
            Dense(1, activation='relu')
        ])

model_compfit(flipann, X_train_flip, y_train, 10)
```

```
Epoch 1/10
233/233 [=====] - 1s 5ms/step - loss: 515.4820 - root_
mean_squared_error: 22.7042
Epoch 2/10
233/233 [=====] - 1s 5ms/step - loss: 491.3006 - root_
mean_squared_error: 22.1653
Epoch 3/10
233/233 [=====] - 1s 5ms/step - loss: 484.7204 - root_
mean_squared_error: 22.0164
Epoch 4/10
233/233 [=====] - 1s 5ms/step - loss: 485.3916 - root_
mean_squared_error: 22.0316
Epoch 5/10
233/233 [=====] - 1s 5ms/step - loss: 487.7504 - root_
mean_squared_error: 22.0851
Epoch 6/10
233/233 [=====] - 1s 5ms/step - loss: 493.7435 - root_
mean_squared_error: 22.2203
Epoch 7/10
233/233 [=====] - 1s 5ms/step - loss: 462.7820 - root_
mean_squared_error: 21.5124
Epoch 8/10
233/233 [=====] - 1s 5ms/step - loss: 467.0221 - root_
mean_squared_error: 21.6107
Epoch 9/10
233/233 [=====] - 1s 5ms/step - loss: 464.9805 - root_
mean_squared_error: 21.5634
Epoch 10/10
233/233 [=====] - 1s 5ms/step - loss: 472.2879 - root_
mean_squared_error: 21.7322
Model: "sequential_2"
```

Layer (type)	Output Shape	Param #
flatten_2 (Flatten)	(None, 49152)	0
dense_2 (Dense)	(None, 1)	49153

Total params: 49,153
 Trainable params: 49,153
 Non-trainable params: 0

```
In [27]: model_eval(flipann, X_train_flip, y_train, X_test, y_test, 'Flipped ANN')
```

```
233/233 [=====] - 2s 8ms/step - loss: 429.9602 - root_mean_squared_error: 20.7355
```

```
78/78 [=====] - 1s 7ms/step - loss: 502.3939 - root_mean_squared_error: 22.4141
```

Flipped ANN Training Metrics:

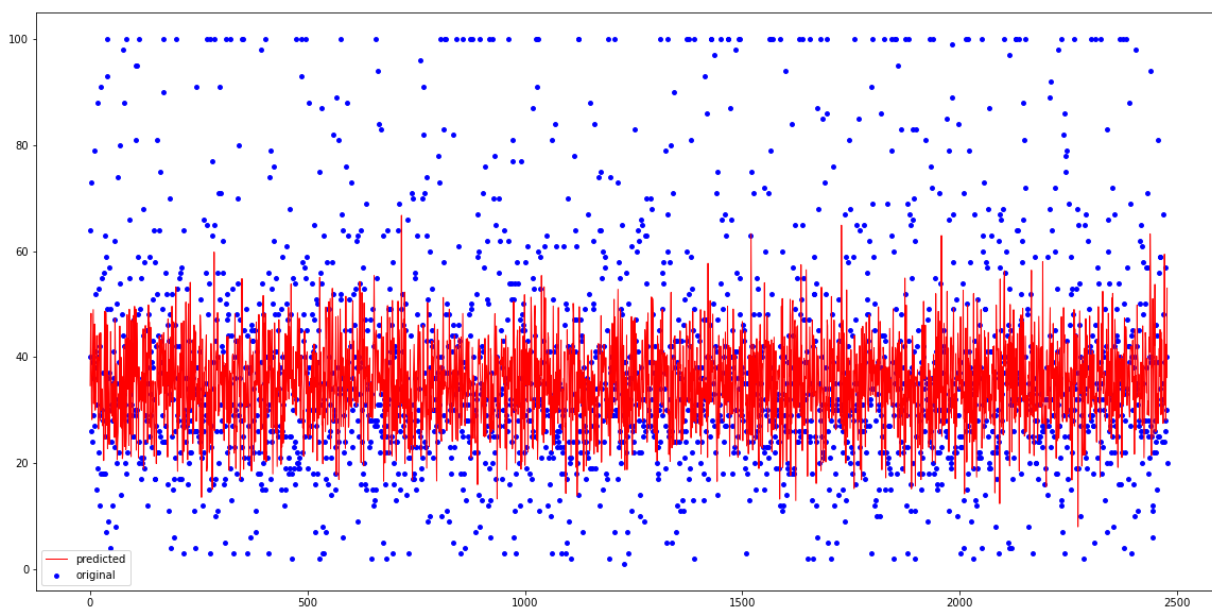
Loss: 429.96

RMSE: 20.735

Flipped ANN Test Metrics:

Loss: 502.394

RMSE: 22.414



Both types of image augmentation performed worse than our baseline.

Iterative Artificial Neural Networks

We move on to build more complex ANN models.

```
In [28]: ann2 = Sequential([
            Flatten(),
            Dense(1000, activation='relu'),
            Dense(1, activation='relu')
        ])

model_compfit(ann2, X_train, y_train, 10)
```

```
Epoch 1/10
233/233 [=====] - 87s 372ms/step - loss: 2517.0427 - r
oot_mean_squared_error: 50.1701
Epoch 2/10
233/233 [=====] - 87s 373ms/step - loss: 1861.6760 - r
oot_mean_squared_error: 43.1471
Epoch 3/10
233/233 [=====] - 81s 348ms/step - loss: 1861.6760 - r
oot_mean_squared_error: 43.1471
Epoch 4/10
233/233 [=====] - 75s 320ms/step - loss: 1861.6760 - r
oot_mean_squared_error: 43.1471
Epoch 5/10
233/233 [=====] - 72s 308ms/step - loss: 1861.6760 - r
oot_mean_squared_error: 43.1471
Epoch 6/10
233/233 [=====] - 72s 310ms/step - loss: 1861.6760 - r
oot_mean_squared_error: 43.1471
Epoch 7/10
233/233 [=====] - 72s 308ms/step - loss: 1861.6760 - r
oot_mean_squared_error: 43.1471
Epoch 8/10
233/233 [=====] - 72s 308ms/step - loss: 1861.6760 - r
oot_mean_squared_error: 43.1471
Epoch 9/10
233/233 [=====] - 72s 311ms/step - loss: 1861.6760 - r
oot_mean_squared_error: 43.1471
Epoch 10/10
233/233 [=====] - 71s 306ms/step - loss: 1861.6760 - r
oot_mean_squared_error: 43.1471
Model: "sequential_3"
```

Layer (type)	Output Shape	Param #
=====		
flatten_3 (Flatten)	(None, 49152)	0
dense_3 (Dense)	(None, 1000)	49153000
dense_4 (Dense)	(None, 1)	1001
=====		
Total params: 49,154,001		
Trainable params: 49,154,001		
Non-trainable params: 0		

```
In [29]: model_eval(ann2, X_train, y_train, X_test, y_test, 'Second ANN')
```

```
233/233 [=====] - 23s 97ms/step - loss: 1861.6760 - root_mean_squared_error: 43.1471  
78/78 [=====] - 7s 96ms/step - loss: 1898.7963 - root_mean_squared_error: 43.5752
```

Second ANN Training Metrics:

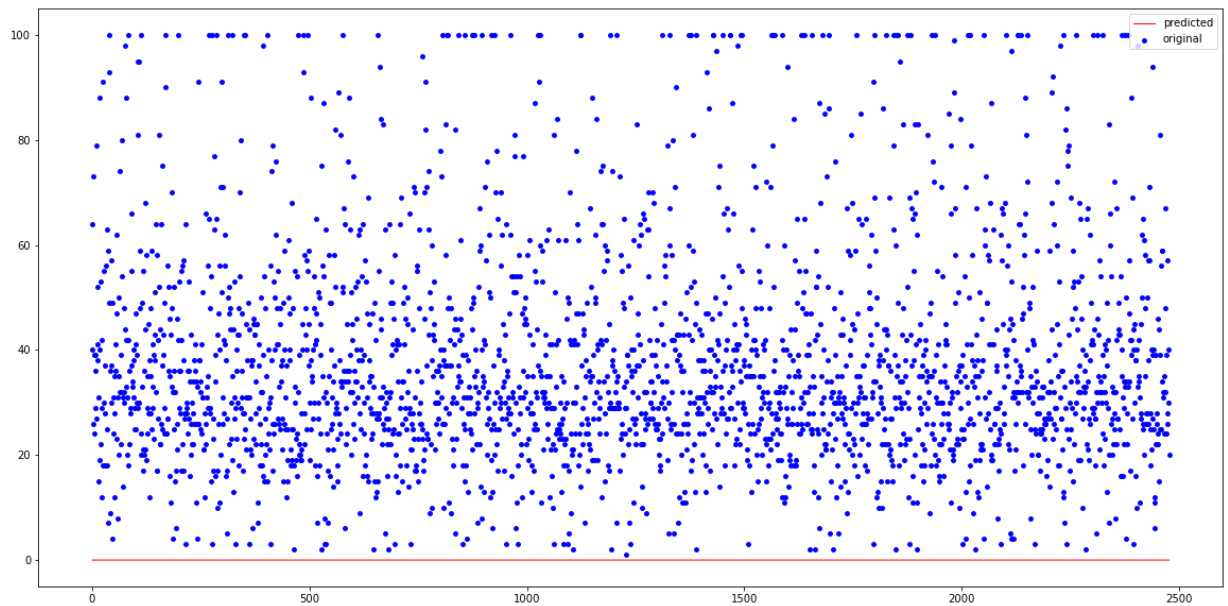
Loss: 1861.676

RMSE: 43.147

Second ANN Test Metrics:

Loss: 1898.796

RMSE: 43.575



```
In [30]: ann3 = Sequential([
            Flatten(),
            Dense(3000, activation='relu'),
            Dense(1000, activation='relu'),
            Dense(10, activation='relu')
        ])

model_compfit(ann3, X_train, y_train, 10)
```

```
Epoch 1/10
233/233 [=====] - 207s 890ms/step - loss: 1239.5125 -
root_mean_squared_error: 35.2067
Epoch 2/10
233/233 [=====] - 193s 830ms/step - loss: 1027.4346 -
root_mean_squared_error: 32.0536
Epoch 3/10
233/233 [=====] - 199s 853ms/step - loss: 1026.6284 -
root_mean_squared_error: 32.0410
Epoch 4/10
233/233 [=====] - 177s 759ms/step - loss: 1021.8079 -
root_mean_squared_error: 31.9657
Epoch 5/10
233/233 [=====] - 166s 713ms/step - loss: 1017.9619 -
root_mean_squared_error: 31.9055
Epoch 6/10
233/233 [=====] - 170s 729ms/step - loss: 1017.0810 -
root_mean_squared_error: 31.8917
Epoch 7/10
233/233 [=====] - 166s 714ms/step - loss: 1013.6356 -
root_mean_squared_error: 31.8376
Epoch 8/10
233/233 [=====] - 166s 714ms/step - loss: 1012.2665 -
root_mean_squared_error: 31.8161
Epoch 9/10
233/233 [=====] - 166s 713ms/step - loss: 1008.4465 -
root_mean_squared_error: 31.7561
Epoch 10/10
233/233 [=====] - 161s 690ms/step - loss: 1005.7928 -
root_mean_squared_error: 31.7142
Model: "sequential_4"
```

Layer (type)	Output Shape	Param #
=====		
flatten_4 (Flatten)	(None, 49152)	0
dense_5 (Dense)	(None, 3000)	147459000
dense_6 (Dense)	(None, 1000)	3001000
dense_7 (Dense)	(None, 10)	10010
=====		
Total params: 150,470,010		
Trainable params: 150,470,010		
Non-trainable params: 0		

```
In [31]: model_eval(ann3, X_train, y_train, X_test, y_test, 'Third ANN')
```

```
233/233 [=====] - 32s 138ms/step - loss: 995.3323 - root_mean_squared_error: 31.5489
```

```
78/78 [=====] - 11s 138ms/step - loss: 1049.9214 - root_mean_squared_error: 32.4025
```

Third ANN Training Metrics:

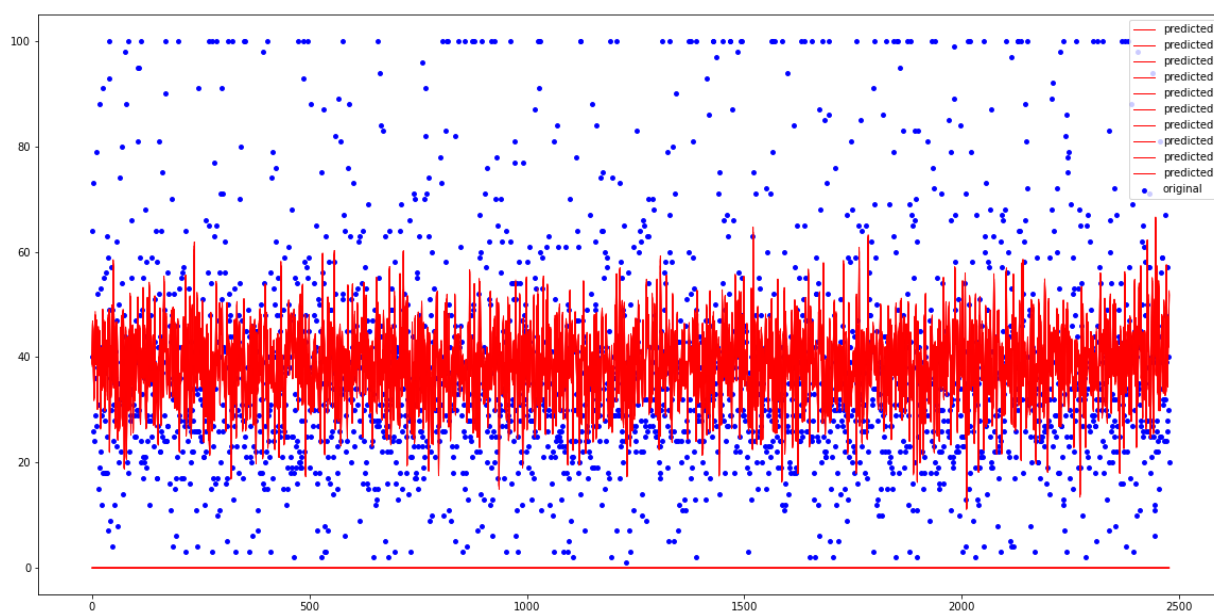
Loss: 995.332

RMSE: 31.549

Third ANN Test Metrics:

Loss: 1049.921

RMSE: 32.402



```
In [32]: ann4 = Sequential([
            Flatten(),
            Dense(units=128, activation='relu'),
            Dense(units=1, activation='relu')
        ])

model_compfit(ann4, X_train, y_train, 10)
```

```
Epoch 1/10
233/233 [=====] - 9s 39ms/step - loss: 1889.0084 - root_mean_squared_error: 43.4627
Epoch 2/10
233/233 [=====] - 8s 33ms/step - loss: 1861.6760 - root_mean_squared_error: 43.1471
Epoch 3/10
233/233 [=====] - 7s 31ms/step - loss: 1861.6760 - root_mean_squared_error: 43.1471
Epoch 4/10
233/233 [=====] - 7s 32ms/step - loss: 1861.6760 - root_mean_squared_error: 43.1471
Epoch 5/10
233/233 [=====] - 7s 32ms/step - loss: 1861.6760 - root_mean_squared_error: 43.1471
Epoch 6/10
233/233 [=====] - 8s 33ms/step - loss: 1861.6760 - root_mean_squared_error: 43.1471
Epoch 7/10
233/233 [=====] - 8s 32ms/step - loss: 1861.6760 - root_mean_squared_error: 43.1471
Epoch 8/10
233/233 [=====] - 7s 31ms/step - loss: 1861.6760 - root_mean_squared_error: 43.1471
Epoch 9/10
233/233 [=====] - 7s 30ms/step - loss: 1861.6760 - root_mean_squared_error: 43.1471
Epoch 10/10
233/233 [=====] - 7s 32ms/step - loss: 1861.6760 - root_mean_squared_error: 43.1471
Model: "sequential_5"
```

Layer (type)	Output Shape	Param #
=====		
flatten_5 (Flatten)	(None, 49152)	0
dense_8 (Dense)	(None, 128)	6291584
dense_9 (Dense)	(None, 1)	129
=====		
Total params: 6,291,713		
Trainable params: 6,291,713		
Non-trainable params: 0		


```
In [33]: model_eval(ann4, X_train, y_train, X_test, y_test, 'Fourth ANN')
```

```
233/233 [=====] - 4s 16ms/step - loss: 1861.6760 - root_mean_squared_error: 43.1471
```

```
78/78 [=====] - 2s 22ms/step - loss: 1898.7963 - root_mean_squared_error: 43.5752
```

Fourth ANN Training Metrics:

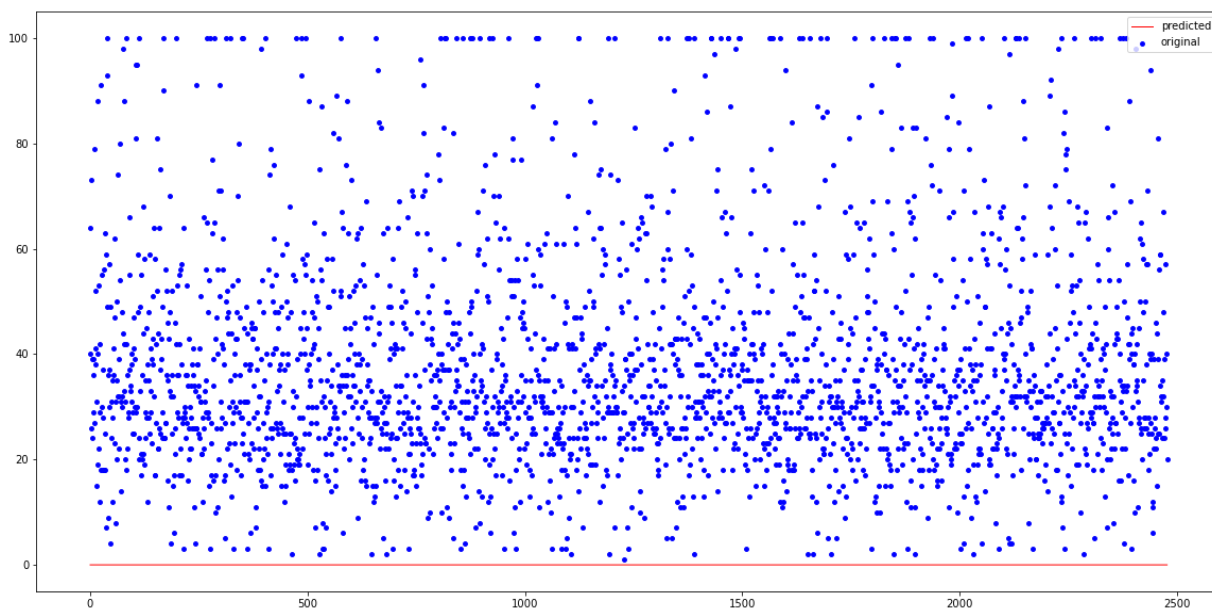
Loss: 1861.676

RMSE: 43.147

Fourth ANN Test Metrics:

Loss: 1898.796

RMSE: 43.575



```
In [35]: ann5 = tf.keras.Sequential([
        Flatten(),
        Dense(units=128, activation='relu'),
        Dense(units=64, activation='relu'),
        Dense(units=1, activation='relu')
    ])

model_compfit(ann5, X_train, y_train, 10)
```

```
Epoch 1/10
233/233 [=====] - 11s 48ms/step - loss: 516.1583 - root_mean_squared_error: 22.7191
Epoch 2/10
233/233 [=====] - 13s 58ms/step - loss: 482.3322 - root_mean_squared_error: 21.9621
Epoch 3/10
233/233 [=====] - 23s 101ms/step - loss: 469.8159 - root_mean_squared_error: 21.6752
Epoch 4/10
233/233 [=====] - 17s 74ms/step - loss: 473.5893 - root_mean_squared_error: 21.7621
Epoch 5/10
233/233 [=====] - 14s 59ms/step - loss: 467.6059 - root_mean_squared_error: 21.6242
Epoch 6/10
233/233 [=====] - 18s 76ms/step - loss: 452.8302 - root_mean_squared_error: 21.2798
Epoch 7/10
233/233 [=====] - 13s 56ms/step - loss: 456.5316 - root_mean_squared_error: 21.3666
Epoch 8/10
233/233 [=====] - 12s 50ms/step - loss: 447.3635 - root_mean_squared_error: 21.1510
Epoch 9/10
233/233 [=====] - 12s 51ms/step - loss: 445.7745 - root_mean_squared_error: 21.1134
Epoch 10/10
233/233 [=====] - 12s 50ms/step - loss: 442.9835 - root_mean_squared_error: 21.0472
Model: "sequential_7"
```

Layer (type)	Output Shape	Param #
flatten_7 (Flatten)	(None, 49152)	0
dense_13 (Dense)	(None, 128)	6291584
dense_14 (Dense)	(None, 64)	8256
dense_15 (Dense)	(None, 1)	65
Total params: 6,299,905		
Trainable params: 6,299,905		
Non-trainable params: 0		

```
In [36]: model_eval(ann5, X_train, y_train, X_test, y_test, 'Fifth ANN')
```

```
233/233 [=====] - 8s 36ms/step - loss: 429.5197 - root_mean_squared_error: 20.7249  
78/78 [=====] - 1s 19ms/step - loss: 487.1628 - root_mean_squared_error: 22.0718
```

Fifth ANN Training Metrics:

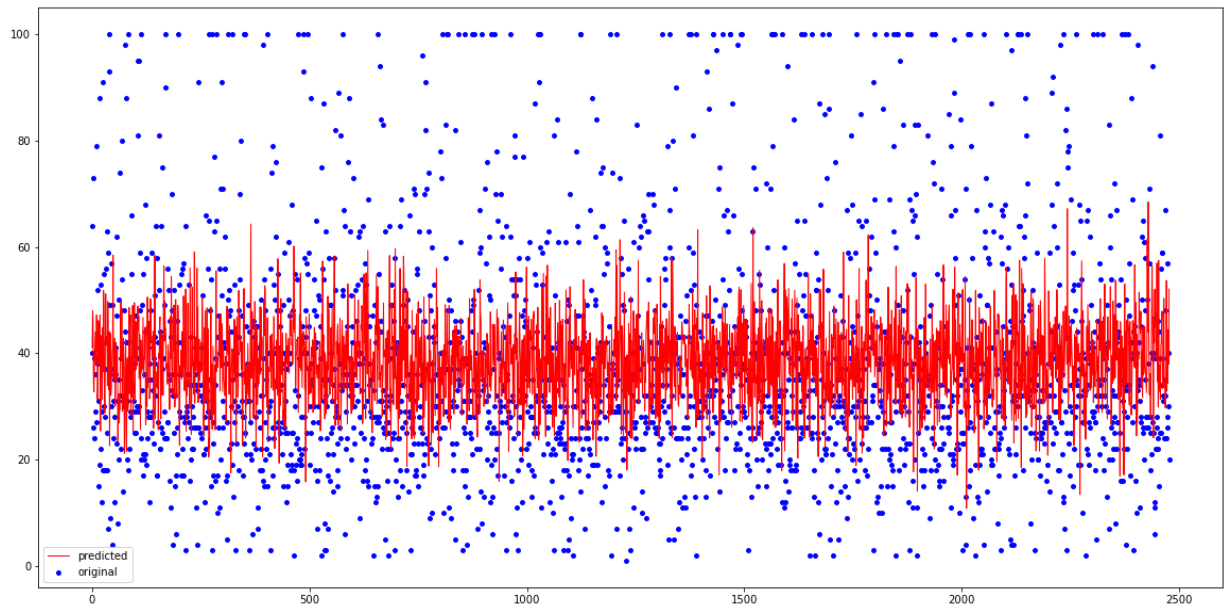
Loss: 429.52

RMSE: 20.725

Fifth ANN Test Metrics:

Loss: 487.163

RMSE: 22.072



We made five total ANN models, and they all performed worse than the baseline.

Iterative Convolutional Neural Networks

We move on to making Convolutional Neural Network models, hoping to do better than our baseline.

```
In [37]: CNN = Sequential([
    Conv2D(32, (3, 3), activation='relu'),
    MaxPooling2D((2, 2)),
    Conv2D(32, (4, 4), activation='relu'),
    MaxPooling2D((2, 2)),
    Conv2D(64, (3, 3), activation='relu'),
    MaxPooling2D((2, 2)),
    Flatten(),
    Dense(128, activation='relu'),
    Dense(64, activation='relu'),
    Dense(1, activation='relu'),
])

model_compfit(CNN, X_train, y_train, 10)
```

```
Epoch 1/10
233/233 [=====] - 202s 869ms/step - loss: 487.6511 -
root_mean_squared_error: 22.0828
Epoch 2/10
233/233 [=====] - 159s 683ms/step - loss: 432.6213 -
root_mean_squared_error: 20.7995
Epoch 3/10
233/233 [=====] - 159s 680ms/step - loss: 428.8285 -
root_mean_squared_error: 20.7082
Epoch 4/10
233/233 [=====] - 154s 659ms/step - loss: 427.0878 -
root_mean_squared_error: 20.6661
Epoch 5/10
233/233 [=====] - 160s 686ms/step - loss: 424.5081 -
root_mean_squared_error: 20.6036
Epoch 6/10
233/233 [=====] - 211s 906ms/step - loss: 423.5745 -
root_mean_squared_error: 20.5809
Epoch 7/10
233/233 [=====] - 153s 656ms/step - loss: 420.7753 -
root_mean_squared_error: 20.5128
Epoch 8/10
233/233 [=====] - 151s 646ms/step - loss: 422.0416 -
root_mean_squared_error: 20.5437
Epoch 9/10
233/233 [=====] - 153s 657ms/step - loss: 421.1585 -
root_mean_squared_error: 20.5221
Epoch 10/10
233/233 [=====] - 147s 632ms/step - loss: 419.6141 -
root_mean_squared_error: 20.4845
Model: "sequential_8"
```

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 126, 126, 32)	896
max_pooling2d (MaxPooling2D)	(None, 63, 63, 32)	0
conv2d_1 (Conv2D)	(None, 60, 60, 32)	16416
max_pooling2d_1 (MaxPooling2D)	(None, 30, 30, 32)	0

1/7/22, 8:17 AM

finalnotebook - Jupyter Notebook

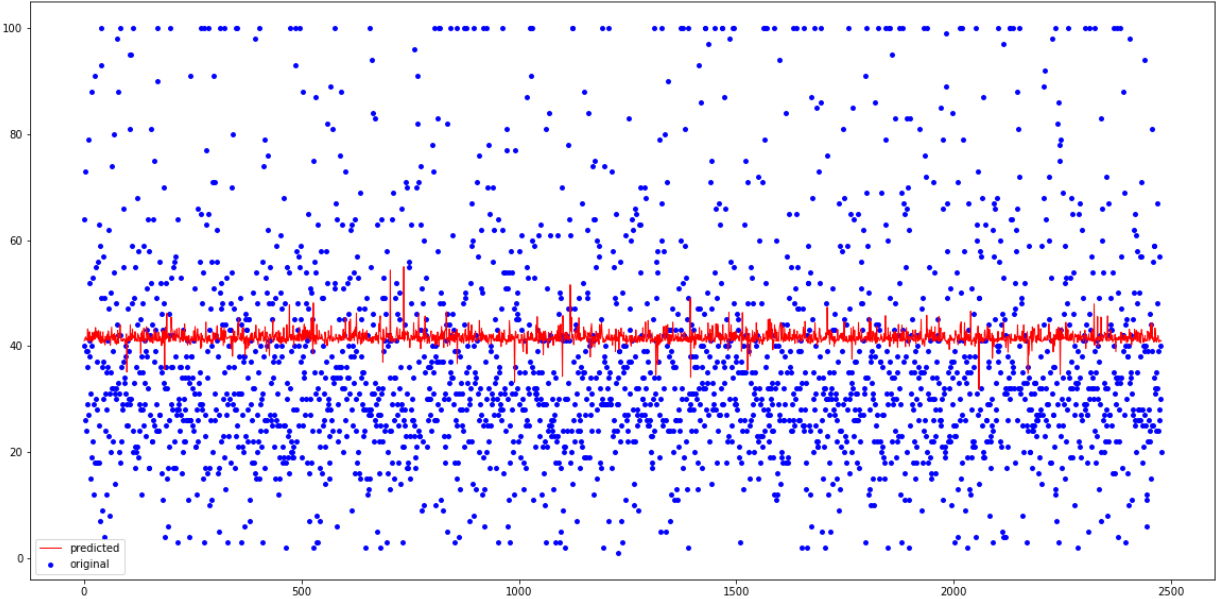
conv2d_2 (Conv2D)	(None, 28, 28, 64)	18496
max_pooling2d_2 (MaxPooling2	(None, 14, 14, 64)	0
flatten_8 (Flatten)	(None, 12544)	0
dense_16 (Dense)	(None, 128)	1605760
dense_17 (Dense)	(None, 64)	8256
dense_18 (Dense)	(None, 1)	65
=====		
Total params: 1,649,889		
Trainable params: 1,649,889		
Non-trainable params: 0		

In [44]: model_eval(CNN, X_train, y_train, X_test, y_test, 'First CNN')

233/233 [=====] - 30s 127ms/step - loss: 425.6042 - ro
ot_mean_squared_error: 20.6302
78/78 [=====] - 11s 142ms/step - loss: 459.9646 - root
_mean_squared_error: 21.4468

First CNN Training Metrics:
Loss: 425.604
RMSE: 20.63

First CNN Test Metrics:
Loss: 459.965
RMSE: 21.447



```
In [45]: CNN2 = Sequential([
    Conv2D(32, (3, 3), activation='relu'),
    MaxPooling2D((2, 2)),
    Flatten(),
    Dense(128, activation='relu'),
    Dense(64, activation='relu'),
    Dense(1, activation='relu'),
])

model_compfit(CNN2, X_train, y_train, 10)
```

Epoch 1/10

233/233 [=====] - 87s 374ms/step - loss: 479.8156 - root_mean_squared_error: 21.9047

Epoch 2/10

233/233 [=====] - 83s 355ms/step - loss: 434.9066 - root_mean_squared_error: 20.8544

Epoch 3/10

233/233 [=====] - 79s 337ms/step - loss: 424.2919 - root_mean_squared_error: 20.5983

Epoch 4/10

233/233 [=====] - 80s 343ms/step - loss: 403.6696 - root_mean_squared_error: 20.0915

Epoch 5/10

233/233 [=====] - 80s 344ms/step - loss: 379.1942 - root_mean_squared_error: 19.4729

Epoch 6/10

233/233 [=====] - 73s 314ms/step - loss: 329.8124 - root_mean_squared_error: 18.1607

Epoch 7/10

233/233 [=====] - 73s 313ms/step - loss: 266.5419 - root_mean_squared_error: 16.3261

Epoch 8/10

233/233 [=====] - 75s 320ms/step - loss: 214.0641 - root_mean_squared_error: 14.6309

Epoch 9/10

233/233 [=====] - 74s 317ms/step - loss: 161.1557 - root_mean_squared_error: 12.6947

Epoch 10/10

233/233 [=====] - 73s 312ms/step - loss: 112.3282 - root_mean_squared_error: 10.5985

Model: "sequential_11"

Layer (type)	Output Shape	Param #
conv2d_5 (Conv2D)	(None, 126, 126, 32)	896
max_pooling2d_6 (MaxPooling2D)	(None, 63, 63, 32)	0
flatten_11 (Flatten)	(None, 127008)	0
dense_25 (Dense)	(None, 128)	16257152
dense_26 (Dense)	(None, 64)	8256
dense_27 (Dense)	(None, 1)	65

```
=====
Total params: 16,266,369
Trainable params: 16,266,369
Non-trainable params: 0
=====
```

```
In [51]: model_eval(CNN2, X_train, y_train, X_test, y_test, 'Second CNN')
```

```
233/233 [=====] - 34s 147ms/step - loss: 74.1564 - root_mean_squared_error: 8.6114
78/78 [=====] - 9s 118ms/step - loss: 602.0927 - root_mean_squared_error: 24.5376
```

Second CNN Training Metrics:

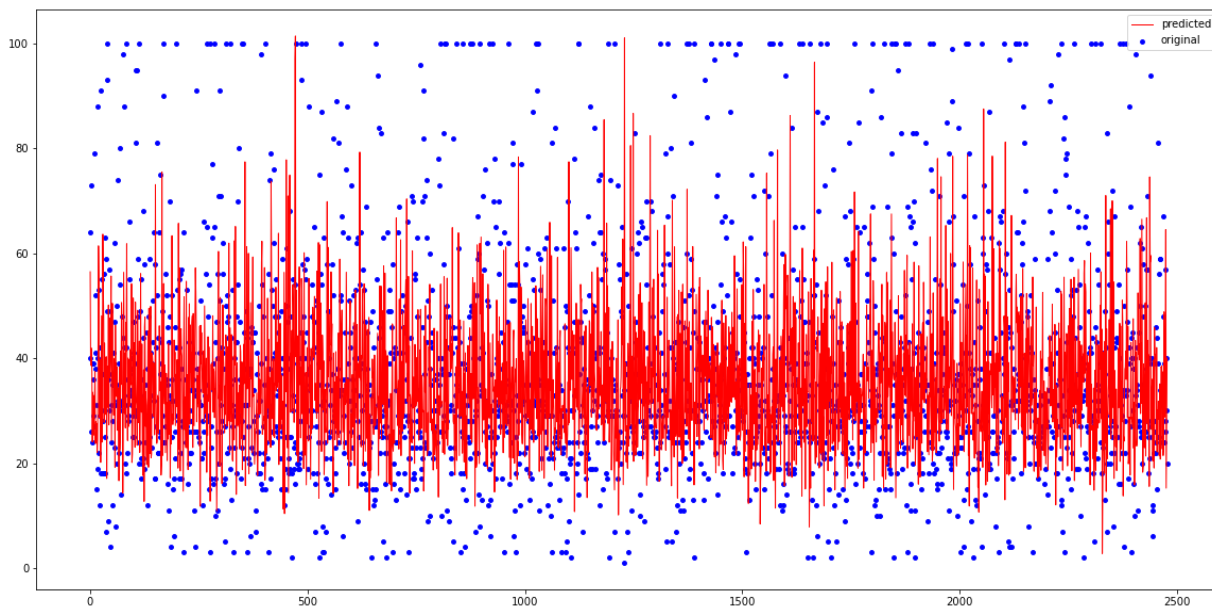
Loss: 74.156

RMSE: 8.611

Second CNN Test Metrics:

Loss: 602.093

RMSE: 24.538



```
In [52]: CNN3 = Sequential([
    Conv2D(32, (3, 3), activation='relu'),
    MaxPooling2D((2, 2)),
    Dense(128, activation='relu'),
    Dropout(0.25),
    Dense(64, activation='relu'),
    MaxPooling2D(pool_size=(2, 2)),
    Flatten(),
    Dropout(0.5),
    Dense(1, activation=None),
])

model_compfit(CNN3, X_train, y_train, 10)
```

Epoch 1/10

233/233 [=====] - 194s 835ms/step - loss: 458.8270 - root_mean_squared_error: 21.4202

Epoch 2/10

233/233 [=====] - 167s 718ms/step - loss: 427.9318 - root_mean_squared_error: 20.6865

Epoch 3/10

233/233 [=====] - 155s 666ms/step - loss: 423.3220 - root_mean_squared_error: 20.5748

Epoch 4/10

233/233 [=====] - 156s 667ms/step - loss: 424.7882 - root_mean_squared_error: 20.6104

Epoch 5/10

233/233 [=====] - 156s 669ms/step - loss: 421.2358 - root_mean_squared_error: 20.5240

Epoch 6/10

233/233 [=====] - 158s 680ms/step - loss: 420.5063 - root_mean_squared_error: 20.5063

Epoch 7/10

233/233 [=====] - 157s 675ms/step - loss: 424.1791 - root_mean_squared_error: 20.5956

Epoch 8/10

233/233 [=====] - 156s 668ms/step - loss: 417.7438 - root_mean_squared_error: 20.4388

Epoch 9/10

233/233 [=====] - 153s 655ms/step - loss: 416.7375 - root_mean_squared_error: 20.4142

Epoch 10/10

233/233 [=====] - 159s 681ms/step - loss: 414.6174 - root_mean_squared_error: 20.3622

Model: "sequential_14"

Layer (type)	Output Shape	Param #
=====		
conv2d_8 (Conv2D)	(None, 126, 126, 32)	896
=====		
max_pooling2d_11 (MaxPooling)	(None, 63, 63, 32)	0
=====		
dense_34 (Dense)	(None, 63, 63, 128)	4224
=====		
dropout_6 (Dropout)	(None, 63, 63, 128)	0
=====		
dense_35 (Dense)	(None, 63, 63, 64)	8256

max_pooling2d_12 (MaxPooling)	(None, 31, 31, 64)	0
flatten_14 (Flatten)	(None, 61504)	0
dropout_7 (Dropout)	(None, 61504)	0
dense_36 (Dense)	(None, 1)	61505
=====		
Total params: 74,881		
Trainable params: 74,881		
Non-trainable params: 0		

In [53]: `model_eval(CNN3, X_train, y_train, X_test, y_test, 'Third CNN')`

233/233 [=====] - 47s 201ms/step - loss: 451.8073 - root_mean_squared_error: 21.2558
 78/78 [=====] - 14s 186ms/step - loss: 500.4067 - root_mean_squared_error: 22.3698

Third CNN Training Metrics:

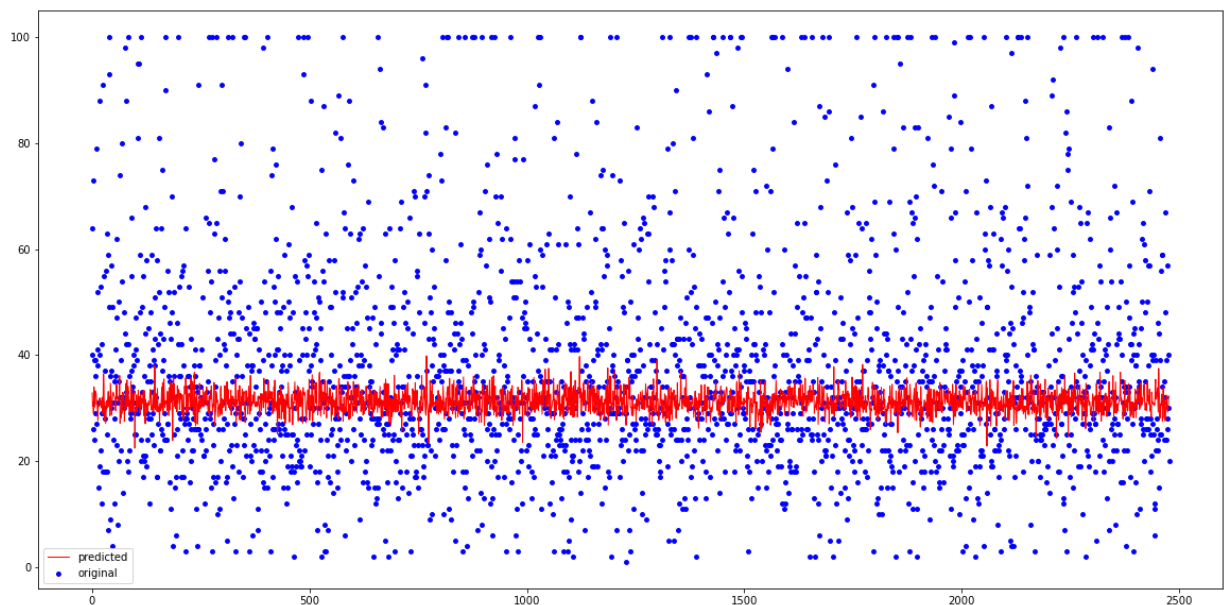
Loss: 451.807

RMSE: 21.256

Third CNN Test Metrics:

Loss: 500.407

RMSE: 22.37



After running these 3 CNN models, we were not able to improve our RMSE from the baseline.

Conclusion

We attempted Data Augmentation but found that Blurring and Flipping the images did not improve our RMSE.

One big concern for us was that how Pawpularity was determined was unclear. There wasn't much information on the Kaggle competition description that explained how the score was created and how to interpret it. We know the scale is from 1-100, but do all the images with a score of 100 have the same amount of traffic? We think understanding how Pawpularity is scored would help us prepare the data for better model results.

As we only had three days to complete this project, this is the best we could do.

Next Steps

First, we would look into what goes into determining Pawpularity. Another thing we would like to look into would be features other than photos that affect pet profile traffic such as age of the animal and time they've been up for adoption. Lastly, our knowledge of neural networks is limited, and so, given more time, we would have liked to explore further.

Sources

- [Kaggle Competition Dataset \(https://www.kaggle.com/c/petfinder-pawpularity-score\)](https://www.kaggle.com/c/petfinder-pawpularity-score)
- [Photo Metadata \(https://github.com/stevenaddison/Project-4/blob/main/data/metadata.md\)](https://github.com/stevenaddison/Project-4/blob/main/data/metadata.md)
- [Speed of Dog Adoption: Impact of Online Photo Traits \(https://www.tandfonline.com/doi/full/10.1080/10888705.2014.982796\)](https://www.tandfonline.com/doi/full/10.1080/10888705.2014.982796)